



LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LIVE WORKSHOP

AI FOUNDATIONS FOR FOUNDERS

Find your highest-value workflow.
Write prompts that work.
Launch a responsible 30-day pilot.

BEGINNER • 2 HOURS • PRACTICAL

AI SHOULD WORK FOR YOU.

AI Foundations for Founders

Understand AI, Choose the Right Work, and Launch Your First Responsible Pilot

WORKSHOP PROMISE

Each learner completes a Founder AI Pilot Blueprint containing a prioritized workflow, tool and data decision, CRAFT prompt, test and verification plan, human-approval controls, success measures, stop conditions, and a 30-day implementation roadmap.

Audience	Lux Automaton founders, operators, small-business leaders, startup teams, and independent professionals
Level	Beginner
Age / audience range	Adults 18+; suitable for mixed-experience business teams
Core duration	120 minutes
Optional clinic	30 minutes
URL slug	ai-foundations-for-founders

Contents

1. Workshop overview
2. Source analysis
3. Learning framework
4. Complete modules and lessons
5. Final project and rubric
6. Visual-production package
7. Publishing package
8. Quality-assurance report
9. Structured JSON companion

1

Workshop overview

A practical two-hour workshop that helps founders understand what AI is, find the best first automation opportunity, protect business data, write useful prompts, and leave with a responsible 30-day pilot plan.

Catalog description

A practical two-hour workshop that helps founders understand what AI is, find the best first automation opportunity, protect business data, write useful prompts, and leave with a responsible 30-day pilot plan.

Detailed workshop overview

AI Foundations for Founders is a beginner-friendly, execution-focused workshop for business owners who want to move beyond AI hype. Learners build a plain-language mental model, audit real operational friction, compare public and private tool environments, write a reusable prompt with the Lux CRAFT framework, and scope one low-risk workflow. Every activity produces part of a final Founder AI Pilot Blueprint. The workshop emphasizes privacy, human approval, measurable outcomes, and reversible experimentation.

Workshop identity

Title	AI Foundations for Founders
Subtitle	Understand AI, Choose the Right Work, and Launch Your First Responsible Pilot
URL slug	ai-foundations-for-founders
Audience	Lux Automaton founders, operators, small-business leaders, startup teams, and independent professionals
Level	Beginner
Age / audience range	Adults 18+; suitable for mixed-experience business teams
Total duration	120-minute core workshop + optional 30-minute implementation clinic
Delivery formats	Live in-person workshop, Live virtual workshop, Facilitated team training, Self-paced workshop with workbook

Final learner outcome

OUTCOME

Each learner completes a Founder AI Pilot Blueprint containing a prioritized workflow, tool and data decision, CRAFT prompt, test and verification plan, human-approval controls, success measures, stop conditions, and a 30-day implementation roadmap.

Materials

- Participant workbook, printed or digital
- Facilitator slide deck
- Laptop or tablet with internet access if an approved AI tool will be used
- A one-paragraph description of the learner's business or business idea
- Sticky notes or a digital whiteboard
- Timer
- Pens or markers
- Synthetic or approved sample business content for prompt testing
- Optional: current tool list, sample SOP, blank customer intake form, or non-sensitive meeting notes

Learning goals

- Understand AI capabilities and limitations in business contexts.
- Identify the highest-value automation opportunity in the learner's business.
- Create a practical AI implementation roadmap.
- Distinguish public, business, and private AI tool environments.
- Write and test a reusable prompt using the Lux CRAFT framework.
- Apply privacy rules, human approval, verification, and stop conditions to an AI pilot.

Prerequisites

- No prior AI experience required.
- Bring a description of your business or idea.
- Bring one recurring workflow or operational frustration.
- Do not bring or display unapproved confidential, regulated, or personally identifiable information.

Safety and privacy notes

REQUIRED SAFETY PRACTICES

- Never share proprietary business data with public AI tools without understanding and approving the privacy implications.
- Use synthetic, public, de-identified, or explicitly approved data during workshop exercises.
- Always verify AI outputs before making or communicating business decisions.
- Keep human oversight on all customer-facing communications and high-impact decisions.
- Do not put passwords, access tokens, authentication codes, payment details, identity documents, or restricted records into general AI prompts.
- Review the terms, data-use commitments, retention, permissions, integrations, and controls for the exact tool plan being used.
- Assign a person with authority to pause or stop the pilot.
- Escalate legal, compliance, privacy, security, employment, financial, or regulated questions to qualified professionals.

Extension activities

- Map three additional workflows for AI assistance or automation.
- Research AI tools specific to the learner's industry using the workshop privacy and control questions.
- Create and run a 30-day AI experiment plan.
- Develop a team AI data rule card and approval process.
- Build a small prompt library with owners, versions, test cases, and review dates.
- Conduct a day-30 evidence review and decide to go, revise, or stop.

Facilitator timing at a glance

#	Module	Time	Learner artifact
1	What AI Actually Is	20 min	AI landscape map
2	Where Automation Helps First	25 min	Opportunity shortlist
3	Private vs. Public AI Tools	25 min	Tool + data rules
4	Writing Useful Prompts	30 min	Tested CRAFT prompt
5	Choosing Your First Workflow	20 min	30-day pilot blueprint



LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

AI FOUNDATIONS FOR FOUNDERS

Understand AI. Find the right workflow. Launch a responsible pilot.

WORKSHOP OUTCOME

Leave with a Founder AI Pilot Blueprint: workflow, tool choice, prompts, guardrails, and a 30-day plan.

5 CORE MODULES • 120 MINUTES

01

WHAT AI ACTUALLY IS

20 min

Plain-language mental model
Capabilities and limits

02

WHERE AUTOMATION HELPS FIRST

25 min

Find friction and repetition
Score value, risk, readiness

03

PRIVATE VS. PUBLIC AI TOOLS

25 min

Classify data first
Choose the right environment

04

WRITING USEFUL PROMPTS

30 min

Use the Lux CRAFT framework
Test, verify, and save

05

CHOOSING YOUR FIRST WORKFLOW

20 min

Scope a low-risk pilot
Build a 30-day roadmap

FOUNDER SAFETY RULES

Protect sensitive data • Verify outputs • Keep humans in charge of customer-facing and high-impact decisions

2

Source analysis

What was preserved, what was expanded, and how factual accuracy is handled.

Supplied source summary

The supplied source defined a five-part, 120-minute beginner curriculum: What AI Actually Is; Where Automation Helps First; Private vs. Public AI Tools; Writing Useful Prompts; and Choosing Your First Workflow. It also specified three learning goals, no-experience prerequisites, privacy and human-oversight safety notes, and three extension activities.

Facts preserved

- The five core topics and their original durations total 120 minutes.
- The audience requires no prior AI experience.
- Learners bring a description of their business or idea.
- The core outcomes are understanding AI capabilities and limitations, identifying the highest-value automation opportunity, and creating an implementation roadmap.
- Sensitive business information should not be placed into public AI tools without understanding privacy implications.
- AI outputs must be verified and customer-facing communications require human oversight.

New curriculum elements

- The five topics are organized as five modules with ten micro-lessons while preserving the original module durations.
- The Lux Opportunity Scorecard provides a repeatable selection method.
- The Lux CRAFT Prompt Framework provides an original prompting structure.
- Every lesson includes a deliverable, case study, knowledge check, facilitator guidance, visual prompt, and alt text.
- The final Founder AI Pilot Blueprint integrates all earlier deliverables.
- New examples are fictional teaching examples and are labeled as such.

Research basis

Definitions and risk-management guidance are aligned with NIST's AI Risk Management Framework and Generative AI Profile. Privacy and confidentiality questions are informed by FTC business guidance. Prompting practices are aligned with official OpenAI guidance while remaining tool-agnostic.

IMPORTANT NOTE

This workshop is educational. It does not replace legal, privacy, cybersecurity, financial, employment, medical, or industry-specific professional advice.

Primary references

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- **Federal Trade Commission (2024). AI Companies: Uphold Your Privacy and Confidentiality Commitments.** <https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2024/01/ai-companies-uphold-your-privacy-confidentiality-commitments>
- **Federal Trade Commission (2026). Artificial Intelligence.** <https://www.ftc.gov/industry/technology/artificial-intelligence>
- **OpenAI (2026). Best practices for prompt engineering with the OpenAI API.** <https://help.openai.com/en/articles/6654000-best-practices-for-prompt-engineering-with-openai-api>
- **OpenAI (2026). Prompt engineering best practices for ChatGPT.** <https://help.openai.com/en/articles/10032626-prompt-engineering-best-practices-for-chatgpt>

3

Learning framework

A five-module path from understanding to implementation.

Instructional approach

The workshop uses an explain-map-build-test-plan sequence. Every concept is immediately connected to a founder decision, every activity creates a visible artifact, and each artifact becomes part of the final Founder AI Pilot Blueprint. The design favors low-risk, reversible experimentation, explicit human approval, and evidence-based decisions.

Learning sequence

Understand	Build a plain-language AI mental model and capability boundary.
Observe	Map actual workflow friction, handoffs, repetition, and exceptions.
Prioritize	Score opportunities by value, readiness, risk, and reversibility.
Protect	Classify data and choose an approved tool environment.
Instruct	Write and test a reusable CRAFT prompt.
Pilot	Scope one workflow and schedule a 30-day evidence-based experiment.

Module structure

MODULE 1: WHAT AI ACTUALLY IS

Build a plain-language mental model of AI, machine learning, generative AI, and automation so founders can make grounded decisions without technical jargon.

Module outcome: Learners can explain the main AI categories, describe where each fits, and state the limits that require human judgment.

Lessons: 1, 2

MODULE 2: WHERE AUTOMATION HELPS FIRST

Turn day-to-day business friction into a prioritized list of AI opportunities based on value, frequency, readiness, risk, and reversibility.

Module outcome: Learners identify and rank the highest-value automation opportunity in their own business.

Lessons: 3, 4

MODULE 3: PRIVATE VS. PUBLIC AI TOOLS

Understand how AI tool categories differ and apply a simple data-classification method before entering information into any system.

Module outcome: Learners choose an appropriate tool class and create clear rules for what business information may be used.

Lessons: 5, 6

MODULE 4: WRITING USEFUL PROMPTS

Use the Lux CRAFT framework to write prompts that are specific, testable, reusable, and aligned with business standards.

Module outcome: Learners create and test a reusable prompt for a real business task.

Lessons: 7, 8


MODULE 5: CHOOSING YOUR FIRST WORKFLOW

Scope a low-risk first workflow and convert it into a 30-day implementation roadmap with owners, metrics, guardrails, and stop conditions.


Module outcome: Learners complete a practical Founder AI Pilot Blueprint ready for controlled testing.

Lessons: 9, 10

Core educational diagrams



The Founder's AI Landscape: Distinguishes rules-based automation, predictive AI, generative AI, and human judgment.




Automation Fit Matrix: Matches repetition and standardization with judgment and consequence to identify human-led, AI-assisted, redesign, or automation candidates.



PUBLIC VS. PRIVATE AI DECISION TREE

Classify the information before choosing the tool.



Public vs. Private AI Decision Tree: Guides tool choice based on the sensitivity of information entering the workflow.



THE LUX CRAFT PROMPT FRAMEWORK

Five elements that turn a vague request into a usable business instruction.




Prompting principle Be clear enough to guide the work, then verify against real standards.

Lux CRAFT Prompt Framework: Explains Context, Role, Assignment, Format plus Boundaries, and Test.

4

Complete modules and lessons

Five modules, ten lessons, and a concrete deliverable from every activity.



MODULE 1 | 20 MINUTES

What AI Actually Is

MODULE DESCRIPTION


Build a plain-language mental model of AI, machine learning, generative AI, and automation so founders can make grounded decisions without technical jargon.

MODULE OUTCOME

Learners can explain the main AI categories, describe where each fits, and state the limits that require human judgment.

Lessons contained in this module

1. Lesson 1: AI, Machine Learning, and Generative AI in Plain Language (10 minutes)
2. Lesson 2: What AI Can Do - and Where It Fails (10 minutes)



LESSON 01 | 10 MINUTES | MODULE 1

AI, Machine Learning, and Generative AI in Plain Language

LESSON OVERVIEW

A founder-friendly explanation of the major AI categories and the input-model-output-human loop behind useful business applications.

Learning objectives

- Define AI in practical business language without treating it as magic or a person.
- Distinguish rules-based automation, predictive AI, and generative AI.
- Map one business task to the type of capability it actually requires.

Teaching lesson

Artificial intelligence is best understood as a family of machine-based capabilities, not as one product or one mind. NIST describes an AI system as a machine-based system that can produce outputs such as predictions, recommendations, or decisions for a defined objective. For founders, the useful question is not “Do I need AI?” It is “What kind of output does this workflow need, and what evidence will make that output trustworthy?”

Start with three categories. Rules-based automation follows explicit instructions: when a form is submitted, create a record, send a confirmation, and assign a task. It is dependable when the steps and exceptions are known. Predictive AI finds patterns in historical examples and estimates a label, score, or likelihood, such as classifying an inquiry by topic or forecasting which inventory item may need attention. Generative AI creates new content from patterns in its training and the context you provide. It can draft text, summarize notes, create variations, extract fields, explain material, or produce an initial plan.

These categories often work together. A founder may use automation to collect a customer form, generative AI to draft a response, and a human to approve the final message. The practical operating model is: input, model or rule, output, review, action. Every step needs an owner. The input must be appropriate for the tool. The output must be checked against the business goal. The action must match the risk.

AI is not a database that always retrieves a stored fact, and fluent wording is not proof of accuracy. Many generative systems produce likely continuations based on patterns. That makes them powerful for drafting and transformation, but it also means they can create plausible errors. AI is also not accountable for the business consequence. The founder or designated operator remains responsible for what is sent, published, purchased, promised, or decided.

A strong founder mental model is therefore simple: AI handles bounded pattern work; people set purpose, provide context, manage exceptions, and approve consequential outcomes. Once learners can name the capability and the human responsibility, they are ready to evaluate where AI belongs in a real workflow.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A boutique marketing agency receives inquiries through a website. Rules-based automation creates a lead record. Generative AI drafts a personalized acknowledgment using the submitted service interest. A human account lead checks claims, tone, and next steps before sending. No AI system decides whether the prospect is a good client without defined criteria and review.

Step-by-step instructions

1. Write one recurring task from your business in a single sentence.
2. Circle the output you need: a fixed action, a prediction or classification, or newly generated content.
3. Label the likely capability: automation, predictive AI, generative AI, or human judgment.
4. Draw the five-part loop: input, capability, output, review, action.
5. Name the person accountable for the final action.

HANDS-ON ACTIVITY

AI Landscape Card Sort: Learners sort eight common founder tasks—such as scheduling, summarizing meeting notes, approving a refund, forecasting demand, drafting a proposal, and updating a CRM—into automation, predictive AI, generative AI, or human-led categories. They may place a task in more than one category if they can explain the sequence.

CONCRETE LEARNER DELIVERABLE

A one-page Founder AI Landscape Map showing one real task, the capability required, the human reviewer, and the final action.

Four useful tips

- Name the business output before naming a tool.
- Separate content creation from business approval.
- Use automation for known steps and AI for pattern-based assistance.
- Treat every AI output as a draft, estimate, or recommendation until verified.

COMMON MISTAKE TO AVOID

Calling every software feature “AI” and skipping the more important question of what the system produces, how it was tested, and who is accountable.

CHECK-IN QUESTION

Which part of your chosen task is pattern work, and which part requires human judgment?

KNOWLEDGE CHECK

What is the clearest difference between rules-based automation and generative AI?

Answer: Rules-based automation follows explicit predefined steps, while generative AI creates new content or transformations from learned patterns and supplied context.

FACILITATOR GUIDANCE

Keep the language concrete. Avoid debating whether AI is “intelligent” in a philosophical sense. Ask learners to anchor every definition to an input, an output, and an accountable person. If participants name a tool, redirect them to the capability and workflow first.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>

Lesson visual and production notes

Image / diagram description	Asa in a premium night technology studio beside a simple visual showing rules-based automation, predictive AI, generative AI, and human judgment connected to a central AI hub.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton workshop visual in a deep navy executive technology studio at night. Show Asa, a confident Black male founder in tailored black clothing and rectangular glasses, on the right. On the left, create a clean, minimal diagram with four large labeled capability cards: Rules-Based Automation, Predictive AI, Generative AI, and Human Judgment, connected to a central AI hub. Use cyan and electric-violet glass lighting, subtle city lights, large readable labels, and clear space for the authentic Lux Automaton logo. Do not generate a logo or tiny interface text.
Accessible alt text	Asa stands in a neon-lit executive studio next to a simple diagram separating automation, predictive AI, generative AI, and human judgment.
Delivered asset	02_lesson_01_what_ai_is.png



LESSON 02 | 10 MINUTES | MODULE 1

What AI Can Do - and Where It Fails

LESSON OVERVIEW

A practical capability-and-limitations map that helps founders use AI for assistance without surrendering verification, judgment, or accountability.

Learning objectives

- Recognize business tasks where AI is strong at transformation, patterning, and first drafts.
- Identify failure modes such as fabricated facts, missing context, inconsistent reasoning, and overconfident language.
- Choose an appropriate level of human review based on consequence and reversibility.

Teaching lesson

AI is most useful when the work is bounded, the desired output is clear, and a person can evaluate the result. Generative AI is often strong at transforming existing material: summarizing a call, turning notes into an outline, rewriting text for a different audience, extracting fields from a consistent document, generating alternatives, or drafting a first version. It can also help a founder compare options, identify questions, and surface patterns that deserve attention.

The same system can fail when the task depends on missing facts, tacit business knowledge, precise legal or financial interpretation, current external information, or consequences that are difficult to

reverse. A polished answer may contain an invented source, an omitted exception, a wrong number, an inappropriate promise, or a confident conclusion that the evidence does not support. The tone of certainty is a writing style; it is not a reliability score.

Use a consequence-based review model. Low-consequence work can receive a quick check. Examples include brainstorming internal meeting names or reorganizing non-sensitive notes. Moderate-consequence work requires verification against a source and a named reviewer. Examples include a proposal draft, a budget explanation, or a customer follow-up. High-consequence work—such as final legal language, credit decisions, employment actions, medical guidance, or commitments that materially affect a person—needs qualified human authority and may be inappropriate for a general AI tool.

A second limit is context. AI cannot automatically know your current pricing, service boundaries, brand promises, internal policies, or customer history unless that information is supplied through an approved system. Even when context is supplied, the model may not prioritize it correctly. That is why founders should provide source material, define boundaries, and specify what the model must not invent.

A third limit is consistency. The same prompt can produce different wording or emphasis across runs. For creative work, variation can help. For operational work, variation must be constrained with templates, examples, checklists, and tests.

The founder rule is: use AI to accelerate thinking and production, not to bypass responsibility. Match the review burden to the possible harm, cost, visibility, and reversibility of the output. If a mistake would be expensive, public, unfair, or difficult to undo, move the human review closer to the decision.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A residential construction consultant uses AI to summarize site notes and draft a client update. The consultant verifies dates, costs, and commitments against the project record. AI may suggest clearer wording, but it never approves a change order or interprets a contract without the responsible professional.

Step-by-step instructions

1. List three tasks AI might help with in your business.
2. For each task, write the possible cost of a wrong output.
3. Mark the output as low, moderate, or high consequence.
4. Choose the review method: quick check, source verification, expert approval, or do not use.
5. Write one sentence describing what the AI must never assume.

HANDS-ON ACTIVITY

Capability Boundary Map: Learners create two columns, “AI can assist” and “Human must decide.” They place parts of one workflow into the appropriate column and add a review rule for every AI-assisted step.

CONCRETE LEARNER DELIVERABLE

A capability boundary map with at least three AI-assisted tasks, their risk level, required reviewer, and one explicit no-go boundary.

Four useful tips

- Ask for options, not unquestioned final answers.
- Require the model to identify missing information.
- Verify facts against the original source, not another AI summary.
- Keep high-impact decisions with qualified people.

COMMON MISTAKE TO AVOID

Using a fluent response as evidence that the answer is correct, complete, current, or appropriate for the business.

CHECK-IN QUESTION

What is one output in your business that looks easy to draft but is expensive to get wrong?

KNOWLEDGE CHECK

Why should review intensity increase with consequence and irreversibility?

Answer: Because errors in high-impact or hard-to-reverse outputs can cause greater harm, cost, or unfairness, so they require stronger human authority and evidence.

FACILITATOR GUIDANCE

Invite participants to name a real error they would not want AI to make. Use those answers to demonstrate consequence-based review. Avoid fear-based framing; the goal is disciplined use, not avoidance. Reinforce that useful AI adoption includes knowing when not to automate.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- **Federal Trade Commission (2026). Artificial Intelligence.** <https://www.ftc.gov/industry/technology/artificial-intelligence>

Lesson visual and production notes

Image / diagram description	LANA in a premium AI studio with a clear split visual: tasks AI can assist with on one side and decisions humans must own on the other.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a deep navy technology studio at night. Show LANA, a warm Black woman AI strategist in a black blazer and lavender blouse, holding a tablet. Beside her, display a clean two-column glass interface titled AI Can Assist and Human Must Decide. Use a few large icons for summarize, draft, classify, approve, commit, and judge. Emphasize human oversight, cyan and electric-violet lighting, subtle city lights, and clear logo placement space. Avoid tiny text and do not generate a fake logo.
Accessible alt text	LANA stands beside a two-column interface separating AI assistance tasks from decisions that require human ownership.
Delivered asset	03_lesson_02_capabilities_limits.png



MODULE 2 | 25 MINUTES

Where Automation Helps First

MODULE DESCRIPTION

Turn day-to-day business friction into a prioritized list of AI opportunities based on value, frequency, readiness, risk, and reversibility.

MODULE OUTCOME

Learners identify and rank the highest-value automation opportunity in their own business.

Lessons contained in this module

1. Lesson 3: Find Repetition, Friction, and Delay (12 minutes)
2. Lesson 4: Score Opportunities by Value, Risk, and Readiness (13 minutes)



LESSON 03 | 12 MINUTES | MODULE 2

Find Repetition, Friction, and Delay

LESSON OVERVIEW

A workflow-discovery lesson that starts with operational pain rather than tools or trends.

Learning objectives

- Identify repeatable work, information bottlenecks, and handoff delays.
- Distinguish a task, a workflow, and a business decision.
- Document one workflow from trigger through outcome and exceptions.

Teaching lesson

Founders often begin AI adoption by browsing tools. A more reliable starting point is the work itself. Look for places where people repeatedly move, rewrite, search for, compare, or wait on information. These are friction signals. They reveal where a better process, automation, or AI assistant may create value.

Common signals include copying the same details between systems, rewriting similar messages, searching across folders for the latest file, converting notes into a standard format, checking whether required information is missing, routing work to the right person, and waiting for an approval that has no visible owner. Not every frustration needs AI. Some problems are caused by unclear ownership,

duplicate tools, or a missing standard operating procedure. The discovery goal is to understand the process before selecting the solution.

Separate three levels. A task is one unit of work, such as drafting a follow-up. A workflow is the connected sequence from a trigger to an outcome, such as receiving a lead, qualifying it, sending a response, scheduling a call, and updating the customer record. A decision is a choice that changes what happens next, such as whether to accept a project or issue a refund. AI may assist a task, automation may move the workflow, and a human may own the decision.

Map one workflow with seven questions: What triggers it? What inputs are required? What steps happen today? Who owns each step? What output is produced? Where does work wait or repeat? What exceptions break the normal path? The exception question matters because a process that looks simple at first may depend on judgment, unavailable data, or customer-specific promises.

Measure friction with evidence that is easy to collect: how often the workflow occurs, approximate minutes spent, number of handoffs, rework, missed deadlines, customer confusion, or time spent searching. Do not invent a precise business case before you have a baseline. A rough, honest estimate is better than a dramatic claim.

The best early opportunity is usually close to existing work, not a full reinvention. It removes a clear burden, uses information the business already has, and keeps a person near the output. By the end of this lesson, learners should have a picture of the current workflow, including the parts that need process improvement before AI is added.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A catering company receives event inquiries by email, website form, and social messages. Staff repeatedly copy dates, guest counts, dietary notes, and venue details into a spreadsheet. The first opportunity is not an autonomous sales agent. It is one structured intake workflow that captures required fields, flags missing information, and prepares a draft response for review.

Step-by-step instructions

1. Choose one business outcome that is frequently delayed or repeated.
2. Name the trigger and final outcome.
3. List the current steps in order, including every handoff.
4. Mark friction with one of four labels: repeat, search, wait, or rework.
5. Record at least one exception and one missing-data problem.
6. Estimate frequency and time using a conservative baseline.

HANDS-ON ACTIVITY

Friction Inventory Sprint: Learners map one workflow on sticky notes or in the workbook, then highlight repeated work in violet, waiting in blue, search in cyan, and rework in pink.

CONCRETE LEARNER DELIVERABLE

A current-state workflow map with trigger, inputs, steps, owners, outcome, friction points, baseline, and exceptions.

Four useful tips

- Observe the real process instead of the ideal process.
- Include manual work that happens outside official systems.
- Ask where information is re-entered or reformatted.
- Fix unclear ownership before adding automation.

COMMON MISTAKE TO AVOID

Choosing a shiny AI feature before documenting the current workflow, data, bottleneck, and exceptions.

CHECK-IN QUESTION

Where does your workflow spend more time waiting or searching than producing value?

KNOWLEDGE CHECK

What is the difference between a task and a workflow?

Answer: A task is one unit of work; a workflow is the connected sequence of tasks, owners, decisions, and handoffs from a trigger to an outcome.

FACILITATOR GUIDANCE

Time-box the mapping. Learners may try to document an entire company. Ask them to choose one outcome and one start point. Push for actual handoffs and exceptions. If a participant cannot estimate time, use frequency, rework, or customer delay as the baseline.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>

Lesson visual and production notes**Image / diagram description**

Asa and LANA review a founder's current process at a design table, with visual markers for repeated work, searching, waiting, and rework.

Detailed image-generation prompt	Create a premium 16:9 Lux Automaton workshop visual showing Asa and LANA helping a founder map a real business workflow. Use a dark executive technology studio with cyan and violet architectural lighting. On the table and transparent wall, show a simple process with highlighted friction markers labeled Repeat, Search, Wait, and Rework. Keep labels large and minimal. Show notebooks, a laptop, and a practical consulting scene. Leave clear space for the authentic logo and avoid fake interface text.
Accessible alt text	Asa and LANA map a business process with four highlighted friction types: repeat, search, wait, and rework.
Delivered asset	04_lesson_03_find_friction.png



LESSON 04 | 13 MINUTES | MODULE 2

Score Opportunities by Value, Risk, and Readiness

LESSON OVERVIEW

A practical scorecard for selecting an early AI opportunity that is useful, testable, and safe enough to learn from.

Learning objectives

- Evaluate AI opportunities across value, frequency, readiness, risk, and reversibility.
- Compare a high-value long-term idea with a lower-risk first pilot.
- Select a shortlist supported by observable workflow evidence.

Teaching lesson

A list of possible AI ideas is not yet a strategy. Founders need a way to compare opportunities without being distracted by novelty. The Lux Opportunity Scorecard uses six questions: How much value could the workflow create? How often does it occur? How repetitive or standardized is it? Is the required data available and usable? How risky is a wrong output? How easy is it to reverse, correct, or stop the workflow?

Score the first four dimensions from one to five, where five is favorable. Score risk from one to five, where five is high risk. Score reversibility from one to five, where five means an error can be quickly detected and undone. The numbers are not scientific precision; they force a structured conversation. A

strong first pilot has meaningful value, regular frequency, enough standardization, available inputs, moderate or low risk, and high reversibility.

Do not automatically choose the biggest possible financial impact. A fully autonomous pricing or hiring system may appear valuable but carry high consequence, hard-to-measure fairness concerns, limited data, and complex exceptions. A smaller workflow—such as preparing a meeting brief from approved notes—may teach the team more quickly and safely. Early adoption should build operational evidence, not merely produce a demonstration.

Add two qualitative questions. First, is a knowledgeable reviewer close to the output? Second, can the workflow be tested in a shadow mode, where AI produces a result but the business does not act on it until the team compares it with current practice? These conditions reduce learning cost and make problems visible.

Rank at least three opportunities. Then write a one-sentence rationale for the top choice and a one-sentence reason the others are not first. This prevents a score from hiding an important concern. A founder may deliberately postpone a high-scoring workflow because a policy, integration, or customer expectation is not ready.

The result should be a shortlist, not a permanent roadmap. New evidence may change the ranking. The scorecard is most valuable when it is connected to the current-state workflow map from the previous lesson. That map provides the frequency, friction, data, and exception evidence needed to score honestly.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A consulting firm considers three ideas: drafting weekly client updates, reviewing final contract language, and generating first-pass research summaries. Contract review has high potential value but high consequence and limited reversibility. Weekly updates have frequent use, clear source notes, a nearby reviewer, and easy correction. The firm chooses weekly updates as the first pilot.

Step-by-step instructions

1. List three possible AI-assisted workflows from your friction inventory.
2. Score value, frequency, standardization, data readiness, risk, and reversibility from one to five.
3. Add whether a knowledgeable reviewer is available and whether shadow testing is possible.
4. Rank the three opportunities.
5. Write a short rationale for the top choice and a delay reason for the others.
6. Confirm that the chosen workflow has a measurable current-state baseline.

HANDS-ON ACTIVITY

Opportunity Ranking Lab: Learners complete the scorecard for three real workflows, then compare their top choice with a partner who challenges one assumption.

CONCRETE LEARNER DELIVERABLE

A scored shortlist of three AI opportunities with evidence, reviewer availability, shadow-test feasibility, and a recommended first candidate.

Four useful tips

- Use conservative scores when evidence is weak.
- Treat high risk as a reason for stronger controls, not a reason to hide it.
- Prefer workflows with a reviewer close to the output.
- Choose a pilot that can be stopped without disrupting the business.

COMMON MISTAKE TO AVOID

Selecting the idea with the most impressive demo instead of the workflow with the best combination of value, readiness, and controllable risk.

CHECK-IN QUESTION

Which score in your top opportunity is based on evidence, and which is still an assumption?

KNOWLEDGE CHECK

Why is reversibility important for a first AI pilot?

Answer: A reversible workflow makes errors easier to detect, correct, and stop, which lowers the cost of learning and protects customers and the business.

FACILITATOR GUIDANCE

Ask learners to defend scores with observable evidence. If everyone gives every idea a five, introduce a forced ranking. Reinforce that the scorecard supports judgment rather than replacing it. Encourage shadow testing for uncertain use cases.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>

Lesson visual and production notes

Image / diagram description	LANA leads a small founder team through a transparent opportunity scorecard with large categories for value, frequency, readiness, risk, and reversibility.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton training scene in a glass-walled executive studio at night. Show LANA presenting a simple transparent scorecard to a diverse founder team. Include five large readable categories: Value, Frequency, Readiness, Risk, and Reversibility, with clear one-to-five indicators. Use deep navy, cyan, and electric-violet light. Keep the composition practical and editorial with negative space for a headline and authentic logo placement. Avoid tiny text.
Accessible alt text	LANA presents a five-part AI opportunity scorecard to a small team in a neon-lit executive studio.
Delivered asset	05_lesson_04_score_opportunities.png



MODULE 3 | 25 MINUTES

Private vs. Public AI Tools

MODULE DESCRIPTION

Understand how AI tool categories differ and apply a simple data-classification method before entering information into any system.

MODULE OUTCOME

Learners choose an appropriate tool class and create clear rules for what business information may be used.

Lessons contained in this module

1. Lesson 5: Public, Private, and Enterprise AI: What Changes? (12 minutes)
2. Lesson 6: Classify Data Before You Prompt (13 minutes)



LESSON 05 | 12 MINUTES | MODULE 3

Public, Private, and Enterprise AI: What Changes?

LESSON OVERVIEW

A tool-selection lesson focused on data flow, retention, access, administration, contractual commitments, and operational control.

Learning objectives

- Describe practical differences among public consumer, enterprise or business, and private AI environments.
- Ask the vendor and policy questions that matter before sensitive use.
- Match a workflow's data and consequence to an appropriate tool class.

Teaching lesson

"Public" and "private" are useful starting labels, but founders should look beyond the label to the actual controls. A public consumer AI tool is typically designed for individual use with standard terms and limited organizational administration. An enterprise or business environment may add contractual data commitments, administrator controls, identity management, retention options, audit features, and separation between organizational users. A private environment may be a dedicated service, a controlled cloud deployment, or a system operated within infrastructure the organization manages.

None of these categories automatically makes a tool safe or unsafe. The correct choice depends on the workflow, the information entering the system, the output, the vendor's terms, the organization's policies, and the controls that are actually configured. The FTC has emphasized that companies should honor privacy and confidentiality promises. Founders should therefore verify data practices instead of relying on marketing language.

Ask seven questions. Is user content used to train or improve shared models? How long are prompts, files, and outputs retained? Who can access them, including vendor staff and subprocessors? Can an administrator manage accounts and permissions? Is data encrypted in transit and at rest? Are audit logs and deletion controls available? What contract, data-processing, or security commitments apply to the specific plan being purchased?

Also ask where integrations send information. A private chat interface connected to an unapproved plug-in, shared drive, or external automation may create a new data path. Tool selection must cover the whole workflow, not just the model window.

Use a simple matching rule. Public tools are best reserved for public, generic, synthetic, or carefully de-identified content when policy permits. Enterprise or approved business tools may support internal or confidential work when the organization has reviewed terms and configured access. Restricted, regulated, credential, or highly sensitive information may require a dedicated private environment, additional legal and security review, or a decision not to use AI.

Finally, remember that output handling matters. An AI-generated document may combine source information in a way that exposes sensitive details. Store, share, and approve outputs according to the same classification as their inputs and intended use.

The goal is not to memorize brand names. It is to build a repeatable tool decision that remains valid as vendors and features change.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A product design studio uses a public tool to brainstorm generic workshop names with fictional examples. It uses an approved business environment for client meeting summaries because the contract and admin controls have been reviewed. It does not upload unreleased product specifications into any tool until the client owner and security lead approve the environment.

Step-by-step instructions

1. Name the workflow and the information that will enter the tool.
2. Classify the information using the next lesson's categories or your existing policy.
3. Review vendor terms for training use, retention, access, deletion, and administration.
4. Check integrations and downstream destinations.
5. Match the workflow to public, approved business, dedicated private, or no-AI use.
6. Record the decision owner and review date.

HANDS-ON ACTIVITY

Tool Class Comparison: Learners compare a public consumer tool, an approved business environment, and a private deployment using the seven control questions. They are not required to select a specific vendor.

CONCRETE LEARNER DELIVERABLE

A tool-class decision table for one workflow, including data type, required controls, approved environment, decision owner, and review date.

Four useful tips

- Verify the terms for the exact plan, not the general brand.
- Review plug-ins, connectors, and automations as part of the data path.
- Use synthetic examples during early experimentation.
- Revisit the decision when the tool, plan, integration, or workflow changes.

COMMON MISTAKE TO AVOID

Assuming the word “private” in marketing guarantees data isolation, limited retention, no training use, or the controls your business requires.

CHECK-IN QUESTION

Which vendor control would change your decision most: training use, retention, access, administration, or auditability?

KNOWLEDGE CHECK

Why is the tool's integration path part of the privacy decision?

Answer: Because connected plug-ins, drives, or automations may send the same information to additional systems with different access, retention, or security controls.

FACILITATOR GUIDANCE

Do not make vendor-specific compliance claims unless the current contract and configuration are known. Teach the questions and decision record. Remind learners that policies and laws vary by location and industry and that legal or security experts may be required.

Recommended resources

- **Federal Trade Commission (2024). AI Companies: Uphold Your Privacy and Confidentiality Commitments.** <https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2024/01/ai-companies-uphold-your-privacy-confidentiality-commitments>

- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).**
<https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- **Federal Trade Commission (2026). Artificial Intelligence.**
<https://www.ftc.gov/industry/technology/artificial-intelligence>

Lesson visual and production notes

Image / diagram description	Asa, LANA, and a business owner review a private-control dashboard showing permissions, files, audit logs, and approvals.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a deep navy executive studio at night. Show Asa, LANA, and a diverse business owner reviewing a transparent dashboard with six large modules: Training Use, Retention, Access, Admin Controls, Audit Logs, and Integrations. Emphasize owner control and privacy. Use cyan and electric-violet glass interfaces, subtle city lights, and clear negative space for the authentic logo. Avoid small illegible copy and do not generate a fake logo.
Accessible alt text	Asa and LANA review privacy, access, retention, and audit controls with a business owner.
Delivered asset	06_lesson_05_public_private.png



LESSON 06 | 13 MINUTES | MODULE 3

Classify Data Before You Prompt

LESSON OVERVIEW

A simple four-level data-classification method and a one-page rule card for safer AI use.

Learning objectives

- Classify information as public, internal, confidential, or restricted.
- Apply data minimization, de-identification, permission, and output-handling rules.
- Create clear green, yellow, and red guidance for a team.

Teaching lesson

Tool selection becomes easier when the business first classifies its information. A simple four-level model works for many small organizations. Public information is approved for anyone to see, such as a published service description. Internal information supports operations but is not intended for public distribution, such as a draft agenda or internal process note. Confidential information could harm the business, a customer, or a partner if disclosed, such as customer history, pricing strategy, non-public financials, or contract terms. Restricted information requires the strongest protection and may include credentials, authentication secrets, regulated records, highly sensitive personal data, or information the business is not authorized to share.

Classification is not only about the prompt. Attached files, conversation history, images, metadata, retrieved knowledge, and generated outputs can all contain sensitive information. A summary of confidential notes remains confidential. A draft that exposes a customer's private situation should not be copied into a public channel simply because AI created it.

Apply four controls. First, minimize: use only the information necessary for the task. Second, de-identify: remove names, account numbers, precise addresses, or unique details when the task does not require them. Third, authorize: confirm that the person using the tool and the tool environment are approved for that data. Fourth, retain deliberately: know where prompts, files, outputs, and logs will be stored and when they should be deleted.

Create green, yellow, and red rules. Green means approved public or synthetic content in an approved tool. Yellow means internal or confidential content that requires an approved business environment, limited access, and human review. Red means do not enter the information without explicit data-owner, legal, security, or compliance approval. Credentials and secrets should remain red.

The rule card must be specific enough for a team member to act without guessing. "Use common sense" is not a control. Better guidance says: "Do not paste customer names, payment details, passwords, contract documents, or unreleased financial data into public AI tools. Use synthetic examples for experiments. For internal summaries, use the approved business account and remove unnecessary identifiers."

This lesson is educational and does not replace legal, security, privacy, or industry-specific advice. The practical objective is to establish a safer default while the organization develops more formal governance.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A recruiting agency wants AI help rewriting job advertisements. Publicly posted job descriptions are green. Internal interview notes are yellow and may only be used in an approved environment with access controls. Candidate identity documents, medical information, credentials, and private demographic data are red and are not entered into a general AI tool.

Step-by-step instructions

1. List the data elements used in your chosen workflow.
2. Classify each as public, internal, confidential, or restricted.
3. Remove any data that is not necessary for the task.
4. Identify whether de-identification is possible.
5. Set the approved tool and reviewer for each class.
6. Write green, yellow, and red examples in plain language.
7. Add an escalation contact for uncertain cases.

HANDS-ON ACTIVITY

Data Rule Card Workshop: Learners create a one-page AI data-handling card for their business with four classifications, examples, allowed environments, prohibited inputs, and escalation instructions.

CONCRETE LEARNER DELIVERABLE

A founder-approved AI data rule card covering input classification, minimization, de-identification, tool approval, output handling, and escalation.

Four useful tips

- Classify the output as carefully as the input.
- Remove identifiers even in an approved tool when they are unnecessary.
- Keep credentials, secrets, and authentication codes out of prompts.
- Give the team a named person to ask when the answer is unclear.

COMMON MISTAKE TO AVOID

Treating a copied excerpt as harmless when the excerpt still reveals a person's identity, a client's confidential context, or a protected business fact.

CHECK-IN QUESTION

What information does your workflow currently collect that the AI task does not actually need?

KNOWLEDGE CHECK

What are the four controls applied after classification?

Answer: Minimize the data, de-identify where possible, confirm authorization, and manage retention deliberately.

FACILITATOR GUIDANCE


Use examples that fit the audience, but do not ask participants to reveal real sensitive data during the workshop. Encourage synthetic examples. State clearly that industry and regional requirements may impose additional obligations.

Recommended resources

- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- **Federal Trade Commission (2024). AI Companies: Uphold Your Privacy and Confidentiality Commitments.** <https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2024/01/ai-companies-uphold-your-privacy-confidentiality-commitments>
- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>

Lesson visual and production notes

Image / diagram description	A transparent data-classification interface organizes public, internal, confidential, and restricted information with clear green, yellow, and red use rules.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton educational visual in a dark executive technology studio. Show LANA and a founder reviewing a clean transparent data-classification board with four large columns: Public, Internal, Confidential, and Restricted. Add a simple green-yellow-red rule strip for Approved, Controlled, and Do Not Enter. Use cyan and electric-violet lighting, minimal readable labels, and space for the authentic logo. Do not include real personal data or fake logos.
Accessible alt text	LANA reviews a four-level data-classification board with public, internal, confidential, and restricted categories.
Delivered asset	07_lesson_06_data_rules.png



MODULE 4 | 30 MINUTES

Writing Useful Prompts

MODULE DESCRIPTION


Use the Lux CRAFT framework to write prompts that are specific, testable, reusable, and aligned with business standards.

MODULE OUTCOME

Learners create and test a reusable prompt for a real business task.

Lessons contained in this module

1. Lesson 7: Write Useful Prompts with the Lux CRAFT Framework (15 minutes)
2. Lesson 8: Test, Refine, Verify, and Save (15 minutes)



LESSON 07 | 15 MINUTES | MODULE 4

Write Useful Prompts with the Lux CRAFT Framework

LESSON OVERVIEW

A five-part prompt framework that makes business instructions clear, bounded, and testable.

Learning objectives

- Use Context, Role, Assignment, Format plus Boundaries, and Test criteria in a prompt.
- Convert a vague request into a specific business instruction.
- Require the model to flag missing information instead of inventing it.

Teaching lesson

A useful prompt is not a clever phrase. It is a compact operating instruction. The Lux CRAFT framework gives founders five elements: Context, Role, Assignment, Format plus boundaries, and Test.

Context explains the situation, audience, source material, and business goal. "Write an email" provides almost no context. "Using the approved call notes below, prepare a follow-up for a prospective bookkeeping client who asked about monthly reporting" gives the model a grounded task.

Role defines the perspective or expertise to apply. A role should guide the work, not pretend the model has real credentials or authority. "Act as an experienced B2B client-communications editor" is useful. "Act as my attorney and make the final legal decision" is not an appropriate transfer of responsibility.

Assignment states the exact task and deliverable. Use an action verb: summarize, compare, extract, draft, classify, or transform. Include the purpose: "Draft a follow-up that confirms the client's priorities and proposes two next steps."

Format plus boundaries defines structure, length, tone, required sections, prohibited claims, and source rules. Tell the model to use only the provided information, to mark unknowns, and not to invent prices, dates, guarantees, testimonials, or commitments. A format may be three short paragraphs, a table, a checklist, JSON, or a reusable template.

Test defines quality. Ask what a good output must accomplish and how it should handle uncertainty. Examples include: every claim must be supported by the notes; all dates must match the source; the tone must be professional and warm; the output must identify missing information; and the final line must contain a clear but non-pressuring next step.

A complete CRAFT prompt can also include examples. One strong example often communicates tone and structure better than several adjectives. Keep examples free of sensitive data unless the environment is approved.

The first output is a test, not the finish. Review it for accuracy, completeness, usefulness, tone, and risk. If the model lacks necessary information, the prompt should instruct it to ask questions or insert a visible placeholder such as "[confirm date]" rather than guessing.

Treat reusable prompts as business assets. Give them a name, owner, version, approved data class, and review date. This turns prompting from improvisation into a repeatable process.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A founder changes "Write a follow-up email" into a CRAFT prompt that includes approved call notes, the role of a client-communications editor, the assignment to confirm three needs and propose a discovery call, a 150-word format, boundaries against invented prices or promises, and tests for source accuracy and a clear next step.

Step-by-step instructions

1. Choose one low- or moderate-consequence drafting or transformation task.
2. Write two sentences of context and identify the approved source material.
3. Define a useful role without transferring real authority.
4. State the exact assignment and business purpose.
5. Specify the format, tone, length, required elements, and prohibited assumptions.
6. Add at least three quality tests and an instruction to flag missing information.
7. Run the prompt only with public, synthetic, or approved data.

HANDS-ON ACTIVITY

CRAFT Prompt Build: Learners turn one vague request into a complete CRAFT prompt and exchange it with a partner who checks whether every element is present.

CONCRETE LEARNER DELIVERABLE

A version-one reusable CRAFT prompt with title, owner, approved data class, source requirements, output format, boundaries, and test criteria.

Four useful tips

- Put the task and source boundaries near the beginning.
- Use examples to show the desired output when consistency matters.
- Tell the model what not to invent.
- Require visible questions or placeholders for missing information.

COMMON MISTAKE TO AVOID

Adding more adjectives instead of adding the missing business context, source material, output structure, and quality criteria.

CHECK-IN QUESTION

Which CRAFT element was missing from the prompt you would normally write?

KNOWLEDGE CHECK

What does the T in CRAFT do?

Answer: Test defines the quality criteria, verification rules, and how the model should handle missing or uncertain information.

FACILITATOR GUIDANCE


Demonstrate a weak prompt and improve it live one element at a time. Keep participants on a real task but require synthetic or approved inputs. Emphasize that roles guide perspective; they do not grant authority or professional accountability.

Recommended resources

- **OpenAI (2026). Best practices for prompt engineering with the OpenAI API.**
<https://help.openai.com/en/articles/6654000-best-practices-for-prompt-engineering-with-openai-api>
- **OpenAI (2026). Prompt engineering best practices for ChatGPT.**
<https://help.openai.com/en/articles/10032626-prompt-engineering-best-practices-for-chatgpt>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).**
<https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>

Lesson visual and production notes

Image / diagram description	LANA presents five glowing CRAFT prompt cards—Context, Role, Assignment, Format plus Boundaries, and Test—in a premium Lux studio.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton workshop visual in a deep navy technology studio at night. Show LANA, a warm Black woman strategist in a black blazer and lavender blouse, pointing to five large glass cards labeled C, R, A, F, T. Expand them with concise labels: Context, Role, Assignment, Format + Boundaries, Test. Use cyan, blue, violet, teal, and magenta accents, large readable text, subtle city lights, and clear space for the authentic logo. Avoid dense copy and fake logos.
Accessible alt text	LANA presents the five-part CRAFT prompt framework on large color-coded glass cards.
Delivered asset	08_lesson_07_craft_prompt.png



LESSON 08 | 15 MINUTES | MODULE 4

Test, Refine, Verify, and Save

LESSON OVERVIEW

A prompt quality loop that converts a first draft into a controlled, reusable business asset.

Learning objectives

- Evaluate AI outputs for accuracy, completeness, usefulness, tone, and risk.
- Improve prompts through controlled iteration and edge-case testing.
- Save a prompt with versioning, ownership, approved data, and a verification checklist.

Teaching lesson

A prompt is only useful if its output survives testing. The Lux prompt loop is Draft, Inspect, Refine, Verify, and Save. Each step produces evidence that the instruction is ready—or shows why it is not.

Draft the first prompt using CRAFT. Then inspect the output against five criteria. Accuracy asks whether claims match the approved source. Completeness asks whether required information and sections are present. Usefulness asks whether the output helps the next person take action. Tone checks whether the wording fits the audience and brand. Risk checks for unsupported claims, sensitive information, unfair assumptions, commitments, or decisions that exceed the model's role.

Refine one important variable at a time. If the output is missing structure, change the format instruction. If it invents details, strengthen the source boundary and missing-information rule. If the tone is

inconsistent, add a short example. Changing everything at once makes it hard to learn which instruction improved the result.

Test normal and edge cases. A normal case contains complete information. An edge case might have a missing date, conflicting notes, an unusual customer request, or a value outside the expected range. The prompt should not hide uncertainty. It should ask a question, flag the conflict, or route the case to a person.

Verification can use a second pass. A verifier prompt asks the model to compare the draft with the source and a checklist, identify unsupported statements, and report uncertainty. This does not replace human review; it helps the reviewer focus. For high-impact content, verification must include the responsible professional or decision owner.

Save the prompt with operational metadata: name, business purpose, owner, version, date, approved input classification, required source, output destination, reviewer, and known limitations. Store a strong example and a failed example. Define when the prompt must be reviewed, such as after a policy, price, product, tool, or workflow change.

Avoid “prompt shopping,” where a user repeatedly asks until the model agrees with a preferred answer. The goal is not a pleasing response. It is a repeatable result that meets real standards. A business prompt is ready for controlled use only when the team can explain what it does, what it does not do, how it is checked, and who can stop it.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A service firm tests a follow-up prompt with three cases: complete notes, a missing project date, and conflicting budget information. The first draft invents a date. The team adds a rule to use “[date needed]” and asks the model to list missing information before drafting. The revised prompt passes all three tests and is saved with a human-review checklist.

Step-by-step instructions

1. Run the CRAFT prompt on one normal synthetic or approved example.
2. Score the output for accuracy, completeness, usefulness, tone, and risk.
3. Change one instruction to address the largest gap.
4. Test one missing-information case and one exception case.
5. Use a verifier prompt to compare the draft with the source and checklist.
6. Document known limitations and mandatory human review.
7. Save the prompt with version, owner, approved data class, and review date.

HANDS-ON ACTIVITY

Prompt Test Lab: Learners run or simulate three test cases, record results, revise one element, and complete a reusable prompt card plus verification checklist.

CONCRETE LEARNER DELIVERABLE

A version-two prompt, a three-case test log, a verifier checklist, and an approval record for controlled use.

Four useful tips

- Change one variable at a time during refinement.
- Include at least one incomplete or conflicting test case.
- Save failed examples because they reveal important boundaries.
- Review the prompt whenever the business source or policy changes.

COMMON MISTAKE TO AVOID

Testing only the ideal example and assuming the prompt will behave safely when information is missing, inconsistent, or unusual.

CHECK-IN QUESTION

What edge case would make your prompt produce a confident but wrong result?

KNOWLEDGE CHECK

Why should prompt changes be made one variable at a time?

Answer: Controlled changes make it easier to identify which instruction caused an improvement or introduced a new problem.

FACILITATOR GUIDANCE

The workshop may not have live access to an AI tool. If not, use sample outputs supplied in the workbook and conduct a paper-based test. Emphasize that a verifier prompt is an aid, not independent proof. Require a named human reviewer.

Recommended resources

- **OpenAI (2026). Best practices for prompt engineering with the OpenAI API.**
<https://help.openai.com/en/articles/6654000-best-practices-for-prompt-engineering-with-openai-api>
- **OpenAI (2026). Prompt engineering best practices for ChatGPT.**
<https://help.openai.com/en/articles/10032626-prompt-engineering-best-practices-for-chatgpt>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).**
<https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>

Lesson visual and production notes

Image / diagram description	Asa and LANA review a prompt test dashboard showing Draft, Inspect, Refine, Verify, and Save with normal and edge-case results.
Detailed image-generation prompt	Create a premium 16:9 Lux Automaton workshop scene in a dark glass office at night. Show Asa and LANA with a founder reviewing a clean transparent prompt-test dashboard. Display five large stages: Draft, Inspect, Refine, Verify, Save. Include three simple test cards: Normal, Missing Information, Exception. Use cyan and violet lighting, readable labels, human approval emphasis, and space for the authentic logo. Avoid tiny text or fake logos.
Accessible alt text	Asa and LANA review a five-stage prompt testing loop with normal, missing-information, and exception cases.
Delivered asset	09_lesson_08_test_prompts.png



MODULE 5 | 20 MINUTES

Choosing Your First Workflow

MODULE DESCRIPTION

Scope a low-risk first workflow and convert it into a 30-day implementation roadmap with owners, metrics, guardrails, and stop conditions.

MODULE OUTCOME

Learners complete a practical Founder AI Pilot Blueprint ready for controlled testing.

Lessons contained in this module

1. Lesson 9: Select and Scope a Low-Risk First Workflow (10 minutes)
2. Lesson 10: Build the 30-Day AI Pilot Roadmap (10 minutes)



LESSON 09 | 10 MINUTES | MODULE 5

Select and Scope a Low-Risk First Workflow

LESSON OVERVIEW

A scoping method that turns the top opportunity into a narrow workflow with clear inputs, outputs, ownership, and exceptions.

Learning objectives

- Select a first workflow using the opportunity scorecard and privacy rules.
- Define a narrow trigger, input, output, reviewer, and exception path.
- Write explicit in-scope, out-of-scope, and stop conditions.

Teaching lesson

A first AI workflow should be small enough to understand and useful enough to matter. The aim is not to prove that AI can do everything. The aim is to learn whether one controlled use case improves a real business outcome.

Start with the top opportunity from the scorecard, then narrow it. Define one trigger, such as “a completed discovery call note is approved.” Define one input set, such as the approved template and call notes. Define one output, such as a draft follow-up email. Name one reviewer. Define where the output goes after approval. This is a pilot boundary.

Write what is in scope and out of scope. In scope might be summarizing confirmed needs, reflecting approved service descriptions, and proposing a next-step meeting. Out of scope might be pricing, contract language, discounts, legal interpretations, promises, or sending without review. These boundaries protect the pilot from quietly becoming a larger system.

Choose a workflow that is reversible and close to human review. The team should be able to compare the AI-assisted result with the existing method, correct errors before they reach a customer, and stop the pilot without disrupting core operations. The workflow should use available, approved data and have enough examples to test normal and exception cases.

Define exceptions before launch. What happens when source notes conflict? What happens when required information is missing? What happens when a customer asks for something outside policy? The safest default is to flag and route, not invent and continue.

Identify the minimum control set: approved tool, approved input class, named reviewer, verification checklist, log of tests or incidents, and a stop owner. For customer-facing work, keep human approval until evidence and policy support a different design. For high-impact decisions, human authority may remain permanent.

Finally, write a one-sentence pilot hypothesis: "If we use an approved AI prompt to draft follow-ups from approved call notes, then account leads can prepare accurate responses faster without increasing correction or complaint rates." The statement should identify the workflow, expected improvement, and guardrail.

A clear scope makes the next lesson possible. Without scope, a 30-day plan becomes a list of tools and hopes. With scope, it becomes an experiment.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A bookkeeping practice chooses new-client follow-up. The pilot begins only after an advisor approves the discovery-call notes. AI drafts a summary and requests missing documents. A human reviews every message. Pricing, tax advice, engagement terms, and automatic sending are explicitly out of scope.

Step-by-step instructions

1. Confirm the top opportunity and its current-state baseline.
2. Write the trigger, approved inputs, desired output, and destination.
3. Name the process owner, prompt owner, and human reviewer.
4. List three in-scope actions and three out-of-scope actions.
5. Define at least three exception routes.
6. Write the minimum control set and stop owner.
7. Create a one-sentence pilot hypothesis.

HANDS-ON ACTIVITY

First Workflow Canvas: Learners complete a one-page canvas with trigger, inputs, AI action, output, reviewer, destination, exceptions, boundaries, controls, and hypothesis.

CONCRETE LEARNER DELIVERABLE

A fully scoped first-workflow canvas suitable for a controlled pilot.

Four useful tips

- Start with one output and one reviewer.
- Keep automatic sending or irreversible action out of the first pilot.
- Route exceptions to a person instead of forcing completion.
- Write out-of-scope items as clearly as in-scope items.

COMMON MISTAKE TO AVOID

Letting the pilot expand from one drafting task into intake, decision-making, pricing, sending, and record changes before the first step is validated.

CHECK-IN QUESTION

What will your pilot refuse to do, even if the AI appears capable?

KNOWLEDGE CHECK

What makes a workflow a strong first pilot?

Answer: It is narrow, valuable, measurable, reversible, based on approved data, close to human review, and easy to stop.

FACILITATOR GUIDANCE

Challenge vague scopes such as “automate marketing” or “build an AI assistant.” Require one trigger and one output. Ask learners to read the out-of-scope section aloud; it often reveals hidden risk.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>

Lesson visual and production notes**Image / diagram description**

Asa and LANA help a founder choose one narrow workflow from a larger operating rhythm and clearly mark the human approval point.

Detailed image-generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a dark executive technology studio at night. Show Asa and LANA with a founder selecting one highlighted workflow from several options. The chosen path should show Trigger, Approved Input, AI Draft, Human Review, and Final Action, with a clear approval gate. Use cyan and electric-violet glass interfaces, simple readable labels, city lights, and authentic logo placement space. Do not generate a logo.
Accessible alt text	Asa and LANA select a five-step first workflow with a visible human approval gate.
Delivered asset	10_lesson_09_choose_workflow.png

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 10 • 10 MIN

BUILD YOUR 30-DAY AI PILOT ROADMAP

Set owners, checkpoints, metrics, guardrails, and a clear go-or-stop decision.

WORKSHOP IMPLEMENTATION ROADMAP
A five-phase path to build your Private AI Business OS.

- 1 AUDIT YOUR TOOLS**
Take inventory of apps, data, and pain points across your business.
- 2 ORGANIZE CORE CONTEXT**
Centralize your files, notes, customers, and business knowledge.
- 3 DESIGN SOP WORKFLOWS**
Turn repeatable tasks into smart, AI-ready SOP workflows.
- 4 SET APPROVALS & PERMISSIONS**
Define who can use, do, and approve what—with full auditability.
- 5 LAUNCH YOUR OPERATING RHYTHM**
Activate daily, weekly, and monthly rhythms that drive momentum.

KEEP DATA PRIVATE
Your data stays under your control. Always private.

KEEP THE HUMAN IN CHARGE
AI supports the work. You make the calls.

KEEP THE NEXT STEP VISIBLE
Clear priorities. Visible progress. Constant forward motion.

LUX AUTOMATON WORKSHOP

LESSON 10 | 10 MINUTES | MODULE 5

Build the 30-Day AI Pilot Roadmap

LESSON OVERVIEW

A practical implementation roadmap with phases, owners, measures, guardrails, review checkpoints, and go-or-stop criteria.

Learning objectives

- Plan a 30-day pilot across baseline, build, test, controlled use, and review phases.
- Choose measures that capture speed, quality, rework, exceptions, and human confidence.
- Set launch gates, incident handling, and go, revise, or stop decisions.

Teaching lesson

A 30-day AI pilot is a structured experiment, not a permanent launch. The roadmap has five phases.

Days 1–5 establish the baseline and controls. Document the current workflow, typical volume, cycle time, rework, customer issues, and quality expectations. Confirm the approved tool, data classification, owner, reviewer, and out-of-scope rules. Select representative test cases without exposing unapproved data.

Days 6–10 build the prompt or automation and the review checklist. Use CRAFT, create a verifier step, and document the workflow. Test internally with normal, incomplete, and exception cases. Record every correction instead of hiding failed outputs.

Days 11–15 run shadow testing. AI produces an output, but the business continues using the current process or holds the result until comparison. Measure whether the AI draft is accurate, complete, usable, and faster to review. Identify recurring failure patterns.

Days 16–25 begin controlled use with a small volume and mandatory human approval. Keep a log of input class, output disposition, reviewer changes, exceptions, incidents, and time saved or added. Do not expand volume because one example looked good.

Days 26–30 review the evidence. Choose one of three decisions: go forward with the same controls, revise and repeat a limited test, or stop. A successful pilot may still be stopped if the tool creates unacceptable privacy, reliability, fairness, customer, or operational risk.

Use a balanced measure set. Speed alone can hide rework. Track cycle time, reviewer effort, acceptance without major edits, error or exception rate, required rework, customer or staff feedback, and incidents. Define success before the pilot begins. Define stop conditions too: unauthorized data use, repeated unsupported claims, missed mandatory steps, material customer harm, inability to audit the workflow, or a failure rate above the team's tolerance.

Assign clear roles: executive sponsor, workflow owner, prompt or configuration owner, reviewer, privacy or security contact, and stop authority. In a small business, one person may hold several roles, but the responsibilities should still be named.

The final Founder AI Pilot Blueprint combines the workflow canvas, tool and data decision, tested prompt, controls, 30-day timeline, measures, and reflection. It is the learner's implementation roadmap and the workshop's final project.

PRACTICAL EXAMPLE / CASE STUDY

Teaching example: A two-person advisory firm pilots AI-assisted meeting summaries. It records the current time required, tests with synthetic notes, runs one week of shadow comparison, then uses AI on five approved meetings with human review. The team stops and revises when action owners are occasionally assigned incorrectly, adds a source-check rule, and repeats the test before wider use.

Step-by-step instructions

1. Write the baseline measures and current workflow volume.
2. Assign owners for the workflow, prompt, review, privacy questions, and stop decision.
3. Schedule the five phases across 30 days.
4. Choose normal, missing-information, and exception test cases.
5. Define success measures and acceptable thresholds.
6. Write stop conditions and incident response steps.
7. Schedule a day-15 checkpoint and day-30 go, revise, or stop review.

HANDS-ON ACTIVITY

30-Day Pilot Builder: Learners place their workflow on a five-phase roadmap, choose measures, assign owners, and complete the final decision criteria.

CONCRETE LEARNER DELIVERABLE

A complete Founder AI Pilot Blueprint and 30-day implementation roadmap.

Four useful tips

- Define success and stop conditions before testing.
- Track reviewer effort, not just AI generation time.
- Keep the first controlled-use volume intentionally small.
- Treat revise or stop as valid, evidence-based outcomes.

COMMON MISTAKE TO AVOID

Calling the pilot successful because the AI produced an impressive example, without comparing it to the baseline or tracking corrections, exceptions, and incidents.

CHECK-IN QUESTION

What evidence would convince you to stop or redesign the pilot, even if the team likes the tool?

KNOWLEDGE CHECK

What are the three valid decisions at the end of the pilot?

Answer: Go forward with controls, revise and repeat a limited test, or stop.

FACILITATOR GUIDANCE

Protect the final ten minutes. Learners need a completed roadmap, not another lecture. Circulate and challenge unclear measures or missing stop conditions. Invite two volunteers to state their pilot hypothesis and final decision criteria.

Recommended resources

- **National Institute of Standards and Technology (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0).** <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- **National Institute of Standards and Technology (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1).** <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf>
- **Federal Trade Commission (2026). Artificial Intelligence.** <https://www.ftc.gov/industry/technology/artificial-intelligence>

Lesson visual and production notes**Image / diagram description**

A five-phase 30-day roadmap from baseline to build, shadow test, controlled use, and evidence review, presented by Asa and LANA.

Detailed image-generation prompt	Create a premium 16:9 Lux Automaton workshop roadmap visual in deep navy, cyan, and electric violet. Show Asa and LANA beside a large five-phase glass timeline labeled Baseline, Build, Shadow Test, Controlled Use, and Review. Add small visible icons for owners, metrics, human approval, incident log, and Go / Revise / Stop. Use large readable labels, subtle city lights, and space for the authentic Lux Automaton logo. Avoid fake logos and tiny interface text.
Accessible alt text	Asa and LANA present a five-phase 30-day AI pilot roadmap ending with go, revise, or stop.
Delivered asset	11_lesson_10_30_day_pilot.png

5

Final project and rubric

Founder AI Pilot Blueprint

Final project brief

Design a controlled, useful, and privacy-aware 30-day pilot for one real business workflow. The blueprint must show why the workflow was selected, what information is allowed, how AI will be instructed and tested, where human judgment remains, how success will be measured, and what evidence will trigger a go, revise, or stop decision.

Required steps

1. Choose one workflow from the opportunity scorecard.
2. Document the current-state trigger, inputs, steps, owners, friction, baseline, and exceptions.
3. Select the tool class and record the data-classification decision.
4. Define in-scope and out-of-scope behavior.
5. Write a CRAFT prompt or automation instruction.
6. Create normal, missing-information, and exception test cases.
7. Define the human reviewer, verification checklist, and approval gate.
8. Set success measures, acceptable thresholds, incidents to log, and stop conditions.
9. Assign owners and schedule the five 30-day phases.
10. Write the final pilot hypothesis and day-30 decision criteria.

Expected deliverables

- Current-state workflow map
- Scored opportunity shortlist
- Tool-class and data-use decision
- AI data rule card
- Version-two CRAFT prompt
- Three-case test log and verifier checklist
- First-workflow canvas
- 30-day pilot roadmap
- Go, revise, or stop criteria

Completion checklist

- ☐ The workflow has one clear trigger and one primary output.
- ☐ The workflow is valuable, measurable, and reversible.
- ☐ The approved input data class is documented.
- ☐ The exact tool environment and decision owner are documented.

- ☐ In-scope and out-of-scope actions are explicit.
- ☐ The prompt includes all five CRAFT elements.
- ☐ Normal, missing-information, and exception cases are tested.
- ☐ A human reviewer and stop authority are named.
- ☐ Measures include quality and rework, not only speed.
- ☐ Stop conditions are written before controlled use.
- ☐ The 30-day phases and checkpoints have owners and dates.

Evaluation rubric

WORKFLOW FIT AND SCOPE

- 4 - Pilot Ready: One clear trigger, output, owner, reviewer, boundaries, baseline, and exception route; the workflow is measurable and reversible.
- 3 - Ready with Minor Revision: The workflow is mostly clear but one boundary, baseline, owner, or exception needs refinement.
- 2 - Developing: The opportunity is useful but scope is broad, hard to measure, or contains several hidden workflows.
- 1 - Beginning: The project is tool-centered, undefined, or not connected to a real business outcome.

PRIVACY, TOOL CHOICE, AND GOVERNANCE

- 4 - Pilot Ready: Data is classified, minimized, and matched to an approved environment; permissions, retention, integrations, escalation, and stop authority are documented.
- 3 - Ready with Minor Revision: The correct general tool class is selected, but one control or decision record needs clarification.
- 2 - Developing: Some privacy concerns are recognized, but the data path or approval rules remain unclear.
- 1 - Beginning: Sensitive information may be used without classification, approval, or control.

PROMPT AND TEST QUALITY

- 4 - Pilot Ready: The prompt uses CRAFT, limits unsupported assumptions, handles missing information, and passes normal and edge-case tests with a verification checklist.
- 3 - Ready with Minor Revision: The prompt is useful but one CRAFT element or test case needs improvement.
- 2 - Developing: The prompt produces a draft but lacks source rules, boundaries, or systematic testing.
- 1 - Beginning: The prompt is vague, untested, or encourages invented information.

HUMAN OVERSIGHT AND RISK CONTROLS

- 4 - Pilot Ready: Human authority, review intensity, exceptions, incident logging, and stop conditions match the consequence of the workflow.
- 3 - Ready with Minor Revision: A reviewer is named but one risk control or exception path is incomplete.
- 2 - Developing: Human review is mentioned but not integrated into the workflow.
- 1 - Beginning: The design relies on autonomous action without appropriate review or stop authority.

30-DAY IMPLEMENTATION AND MEASUREMENT

- 4 - Pilot Ready: The plan includes baseline, build, shadow test, controlled use, review, owners, balanced measures, checkpoints, and go/revise/stop criteria.
- 3 - Ready with Minor Revision: The phases and measures are present but one owner, threshold, or checkpoint is unclear.
- 2 - Developing: The plan lists activities but lacks a baseline, balanced measures, or decision criteria.
- 1 - Beginning: The plan is a general intention without timeline, ownership, or evidence.

SCORING

Maximum 20 points. A pilot-ready submission scores 16 or higher, with no criterion below 3. A lower score indicates where the pilot should be revised before controlled use.

Reflection questions

- What did you assume about AI before the workshop that you would now state more carefully?
- Why is this workflow a better first pilot than the other opportunities you scored?
- What information or customer expectation creates the greatest risk?
- What will the human reviewer look for first?
- What evidence would make you expand the pilot?
- What evidence would make you pause or stop it?
- What must remain a human decision even if the technology improves?

OPTIONAL ADVANCED CHALLENGE

Add a second-stage workflow only after the first pilot reaches its success criteria. Document the handoff between the two workflows, new data permissions, integration risks, monitoring, and a rollback plan.

6

Visual-production package

Delivered visuals, generation direction, and accessibility text.

Delivered asset index

Workshop hero	00_workshop_hero.png
Workshop thumbnail	01_workshop_thumbnail.png
Lesson 1 visual	02_lesson_01_what_ai_is.png
Lesson 2 visual	03_lesson_02_capabilities_limits.png
Lesson 3 visual	04_lesson_03_find_friction.png
Lesson 4 visual	05_lesson_04_score_opportunities.png
Lesson 5 visual	06_lesson_05_public_private.png
Lesson 6 visual	07_lesson_06_data_rules.png
Lesson 7 visual	08_lesson_07_craft_prompt.png
Lesson 8 visual	09_lesson_08_test_prompts.png
Lesson 9 visual	10_lesson_09_choose_workflow.png
Lesson 10 visual	11_lesson_10_30_day_pilot.png
The Founder's AI Landscape	12_diagram_ai_landscape.png
Automation Fit Matrix	13_diagram_automation_fit_matrix.png
Public vs. Private AI Decision Tree	14_diagram_privacy_decision_tree.png
Lux CRAFT Prompt Framework	15_diagram_craft_prompt_framework.png
Printable overview	16_printable_workshop_overview.png
Social promo	17_social_promo_4x5.png

Workshop thumbnail



Generation prompt

Premium Lux Automaton workshop thumbnail in a night executive technology studio. Asa and LANA face the camera with a subtle glass AI workflow behind them. Large clean headline: AI FOUNDATIONS FOR FOUNDERS. Small subline: Find your workflow. Protect your data. Launch a 30-day pilot. Deep navy, cyan, electric violet, city lights, sophisticated editorial composition, authentic logo placement area, no generated logo or tiny interface text.

Alt text

Asa and LANA stand in a neon-lit technology studio beside the title AI Foundations for Founders.

Workshop hero



Generation prompt


Premium Lux Automaton hero image for AI Foundations for Founders. Asa and LANA collaborate at a dark executive desk in a glass-walled studio at night. A clean visual path shows Understand AI, Choose the Right Work, and Launch Your First Pilot. Deep navy, cyan, and electric violet lighting, subtle city skyline, editorial negative space, authentic logo placement area, no fake logo, no dense UI text.

Alt text

Asa and LANA work together in a dark technology studio beneath the workshop promise to understand AI, choose the right workflow, and launch a first pilot.

Lesson visuals

Lesson 1: AI, Machine Learning, and Generative AI in Plain Language



Description	Asa in a premium night technology studio beside a simple visual showing rules-based automation, predictive AI, generative AI, and human judgment connected to a central AI hub.
Generation prompt	Create a premium 16:9 Lux Automaton workshop visual in a deep navy executive technology studio at night. Show Asa, a confident Black male founder in tailored black clothing and rectangular glasses, on the right. On the left, create a clean, minimal diagram with four large labeled capability cards: Rules-Based Automation, Predictive AI, Generative AI, and Human Judgment, connected to a central AI hub. Use cyan and electric-violet glass lighting, subtle city lights, large readable labels, and clear space for the authentic Lux Automaton logo. Do not generate a logo or tiny interface text.
Alt text	Asa stands in a neon-lit executive studio next to a simple diagram separating automation, predictive AI, generative AI, and human judgment.

Lesson 2: What AI Can Do - and Where It Fails



Description	LANA in a premium AI studio with a clear split visual: tasks AI can assist with on one side and decisions humans must own on the other.
Generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a deep navy technology studio at night. Show LANA, a warm Black woman AI strategist in a black blazer and lavender blouse, holding a tablet. Beside her, display a clean two-column glass interface titled AI Can Assist and Human Must Decide. Use a few large icons for summarize, draft, classify, approve, commit, and judge. Emphasize human oversight, cyan and electric-violet lighting, subtle city lights, and clear logo placement space. Avoid tiny text and do not generate a fake logo.
Alt text	LANA stands beside a two-column interface separating AI assistance tasks from decisions that require human ownership.

Lesson 3: Find Repetition, Friction, and Delay



Description	Asa and LANA review a founder's current process at a design table, with visual markers for repeated work, searching, waiting, and rework.
Generation prompt	Create a premium 16:9 Lux Automaton workshop visual showing Asa and LANA helping a founder map a real business workflow. Use a dark executive technology studio with cyan and violet architectural lighting. On the table and transparent wall, show a simple process with highlighted friction markers labeled Repeat, Search, Wait, and Rework. Keep labels large and minimal. Show notebooks, a laptop, and a practical consulting scene. Leave clear space for the authentic logo and avoid fake interface text.
Alt text	Asa and LANA map a business process with four highlighted friction types: repeat, search, wait, and rework.

Lesson 4: Score Opportunities by Value, Risk, and Readiness



Description	LANA leads a small founder team through a transparent opportunity scorecard with large categories for value, frequency, readiness, risk, and reversibility.
Generation prompt	Create a premium 16:9 Lux Automaton training scene in a glass-walled executive studio at night. Show LANA presenting a simple transparent scorecard to a diverse founder team. Include five large readable categories: Value, Frequency, Readiness, Risk, and Reversibility, with clear one-to-five indicators. Use deep navy, cyan, and electric-violet light. Keep the composition practical and editorial with negative space for a headline and authentic logo placement. Avoid tiny text.
Alt text	LANA presents a five-part AI opportunity scorecard to a small team in a neon-lit executive studio.

Lesson 5: Public, Private, and Enterprise AI: What Changes?



Description	Asa, LANA, and a business owner review a private-control dashboard showing permissions, files, audit logs, and approvals.
Generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a deep navy executive studio at night. Show Asa, LANA, and a diverse business owner reviewing a transparent dashboard with six large modules: Training Use, Retention, Access, Admin Controls, Audit Logs, and Integrations. Emphasize owner control and privacy. Use cyan and electric-violet glass interfaces, subtle city lights, and clear negative space for the authentic logo. Avoid small illegible copy and do not generate a fake logo.
Alt text	Asa and LANA review privacy, access, retention, and audit controls with a business owner.

Lesson 6: Classify Data Before You Prompt




Description	A transparent data-classification interface organizes public, internal, confidential, and restricted information with clear green, yellow, and red use rules.
Generation prompt	Create a premium 16:9 Lux Automaton educational visual in a dark executive technology studio. Show LANA and a founder reviewing a clean transparent data-classification board with four large columns: Public, Internal, Confidential, and Restricted. Add a simple green-yellow-red rule strip for Approved, Controlled, and Do Not Enter. Use cyan and electric-violet lighting, minimal readable labels, and space for the authentic logo. Do not include real personal data or fake logos.
Alt text	LANA reviews a four-level data-classification board with public, internal, confidential, and restricted categories.

Lesson 7: Write Useful Prompts with the Lux CRAFT Framework



Description	LANA presents five glowing CRAFT prompt cards—Context, Role, Assignment, Format plus Boundaries, and Test—in a premium Lux studio.
Generation prompt	Create a premium 16:9 Lux Automaton workshop visual in a deep navy technology studio at night. Show LANA, a warm Black woman strategist in a black blazer and lavender blouse, pointing to five large glass cards labeled C, R, A, F, T. Expand them with concise labels: Context, Role, Assignment, Format + Boundaries, Test. Use cyan, blue, violet, teal, and magenta accents, large readable text, subtle city lights, and clear space for the authentic logo. Avoid dense copy and fake logos.
Alt text	LANA presents the five-part CRAFT prompt framework on large color-coded glass cards.

Lesson 8: Test, Refine, Verify, and Save



Description	Asa and LANA review a prompt test dashboard showing Draft, Inspect, Refine, Verify, and Save with normal and edge-case results.
Generation prompt	Create a premium 16:9 Lux Automaton workshop scene in a dark glass office at night. Show Asa and LANA with a founder reviewing a clean transparent prompt-test dashboard. Display five large stages: Draft, Inspect, Refine, Verify, Save. Include three simple test cards: Normal, Missing Information, Exception. Use cyan and violet lighting, readable labels, human approval emphasis, and space for the authentic logo. Avoid tiny text or fake logos.
Alt text	Asa and LANA review a five-stage prompt testing loop with normal, missing-information, and exception cases.

Lesson 9: Select and Scope a Low-Risk First Workflow



Description	Asa and LANA help a founder choose one narrow workflow from a larger operating rhythm and clearly mark the human approval point.
Generation prompt	Create a premium 16:9 Lux Automaton lesson visual in a dark executive technology studio at night. Show Asa and LANA with a founder selecting one highlighted workflow from several options. The chosen path should show Trigger, Approved Input, AI Draft, Human Review, and Final Action, with a clear approval gate. Use cyan and electric-violet glass interfaces, simple readable labels, city lights, and authentic logo placement space. Do not generate a logo.
Alt text	Asa and LANA select a five-step first workflow with a visible human approval gate.


Lesson 10: Build the 30-Day AI Pilot Roadmap



Description	A five-phase 30-day roadmap from baseline to build, shadow test, controlled use, and evidence review, presented by Asa and LANA.
Generation prompt	Create a premium 16:9 Lux Automaton workshop roadmap visual in deep navy, cyan, and electric violet. Show Asa and LANA beside a large five-phase glass timeline labeled Baseline, Build, Shadow Test, Controlled Use, and Review. Add small visible icons for owners, metrics, human approval, incident log, and Go / Revise / Stop. Use large readable labels, subtle city lights, and space for the authentic Lux Automaton logo. Avoid fake logos and tiny interface text.
Alt text	Asa and LANA present a five-phase 30-day AI pilot roadmap ending with go, revise, or stop.

Educational diagrams

The Founder's AI Landscape




PURPOSE

Distinguishes rules-based automation, predictive AI, generative AI, and human judgment.

ALT TEXT

A central AI hub connects to four capability categories: automation, predictive AI, generative AI, and human judgment.

Automation Fit Matrix




PURPOSE

Matches repetition and standardization with judgment and consequence to identify human-led, AI-assisted, redesign, or automation candidates.

ALT TEXT

A four-quadrant matrix places workflows by repetition and required judgment.

Public vs. Private AI Decision Tree




PURPOSE

Guides tool choice based on the sensitivity of information entering the workflow.

ALT TEXT

A decision tree routes non-sensitive information toward approved public tools and sensitive or uncertain information toward approved private environments or no AI use.

Lux CRAFT Prompt Framework



PURPOSE

Explains Context, Role, Assignment, Format plus Boundaries, and Test.

ALT TEXT

Five color-coded cards spell CRAFT and define the five elements of a useful prompt.

Printable workshop overview

**LUX AUTOMATON**
AUTOMATE | INNOVATE | ACCELERATE

AI FOUNDATIONS FOR FOUNDERS

Understand AI. Find the right workflow. Launch a responsible pilot.

WORKSHOP OUTCOME
Leave with a Founder AI Pilot Blueprint: workflow, tool choice, prompts, guardrails, and a 30-day plan.

5 CORE MODULES • 120 MINUTES

01

WHAT AI ACTUALLY IS

20 min

Plain-language mental model
Capabilities and limits

02

WHERE AUTOMATION HELPS FIRST

25 min

Find friction and repetition
Score value, risk, readiness

03

PRIVATE VS. PUBLIC AI TOOLS

25 min

Classify data first
Choose the right environment

04

WRITING USEFUL PROMPTS

30 min

Use the Lux CRAFT framework
Test, verify, and save

05

CHOOSING YOUR FIRST WORKFLOW

20 min

Scope a low-risk pilot
Build a 30-day roadmap

FOUNDER SAFETY RULES

Protect sensitive data • Verify outputs • Keep humans in charge of customer-facing and high-impact decisions

Social-media promotional image

A promotional image for a workshop. It features a man with glasses and a beard, wearing a dark suit, sitting at a desk and looking at a tablet. The background is a blurred cityscape at night. The text is overlaid on the left side of the image.

 **LUX AUTOMATON**
AUTOMATE | INNOVATE | ACCELERATE

LIVE WORKSHOP

AI FOUNDATIONS FOR FOUNDERS

Find your highest-value workflow.
Write prompts that work.
Launch a responsible 30-day pilot.

BEGINNER • 2 HOURS • PRACTICAL

AI SHOULD WORK FOR YOU.

7

Publishing package

Ready-to-use website, email, video, and social copy.

SEO title	AI Foundations for Founders Lux Automaton Workshop
Meta description	Understand AI, find your best first workflow, protect business data, write useful prompts, and build a responsible 30-day AI pilot.
Workshop excerpt	Move from AI curiosity to a controlled business experiment. This practical Lux Automaton workshop helps founders understand capabilities and limits, identify the right first automation opportunity, protect sensitive data, write better prompts, and leave with a complete 30-day pilot blueprint.
Suggested CTA	Bring one recurring business frustration. Leave with one responsible AI workflow ready to test.

Promotional headlines

- Stop Chasing AI Tools. Start With the Right Workflow.
- Build Your First Responsible AI Pilot in Two Hours.
- AI Foundations Every Founder Should Know Before Automating.
- Protect the Data. Improve the Work. Keep the Human in Charge.
- From AI Curiosity to a 30-Day Founder Pilot Blueprint.

Social captions

CAPTION 1

AI should work for your business—not create another layer of confusion. AI Foundations for Founders helps you find the right workflow, protect your data, write useful prompts, and launch a responsible 30-day pilot.

CAPTION 2

Before you automate, learn what AI actually is, where it helps, where it fails, and what must remain human. Join Lux Automaton for a practical two-hour founder workshop.

CAPTION 3

Your first AI win probably is not a fully autonomous agent. It is one clear, measurable, reversible workflow. Build that workflow with us.

CAPTION 4

Public tool or private environment? Draft or decision? Assistance or automation? Learn the founder frameworks that make AI adoption safer and more useful.

CAPTION 5

Leave with more than notes. Leave with a scored opportunity, tested prompt, privacy rules, human approval plan, and 30-day implementation roadmap.

Newsletter announcement

AI is moving quickly, but strong business adoption begins with disciplined fundamentals. Lux Automaton's AI Foundations for Founders is a practical, beginner-friendly workshop for owners and operators who want useful results without surrendering privacy, judgment, or control. In two focused hours, participants learn what AI actually is, identify the best first automation opportunity, compare public and private tool environments, write and test a reusable prompt, and build a responsible 30-day pilot roadmap. Every activity contributes to a final Founder AI Pilot Blueprint that can be reviewed with a team and used to guide a controlled experiment.

YouTube description

AI Foundations for Founders is Lux Automaton's practical introduction to responsible business AI. Learn the difference between automation, predictive AI, and generative AI; identify the workflow that creates the most value; protect sensitive information; use the Lux CRAFT prompt framework; and build a measurable 30-day pilot with human approval and clear stop conditions. This workshop is designed for founders, operators, small-business leaders, and startup teams with no prior AI experience.

Suggested tags

AI for founders, business automation, generative AI, AI strategy, prompt engineering, AI privacy, private AI, small business AI, AI workflow, human in the loop, responsible AI, Lux Automaton

Team-leader announcement

Team training announcement: We are attending AI Foundations for Founders to establish a shared language for AI, identify one appropriate pilot, and create clear privacy and human-approval rules. Please bring one repetitive workflow, one operational bottleneck, and a non-sensitive example that can be used during exercises. Do not bring confidential or regulated data.

8

Quality-assurance report

Verification against the workshop specification.

LESSON COMPLETENESS

Pass - all ten lessons include every required component.

ACTIVITY DELIVERABLES

Pass - every activity produces a named artifact used in the final project.

FINAL PROJECT ALIGNMENT

Pass - the final blueprint combines workflow discovery, scoring, privacy, prompting, testing, oversight, and planning skills taught earlier.

LESSON UNIQUENESS

Pass - no lesson repeats another; each advances a distinct stage of the adoption process.

SAFETY PRIVACY

Pass - the workshop includes data classification, tool controls, human review, escalation, incident logging, and stop authority.

VISUAL ALIGNMENT

Pass - visuals use Lux Automaton's deep navy, cyan, electric-violet, executive-studio direction and use the authentic supplied logo only in post-production.

FACTUAL ACCURACY

Pass - claims are grounded in the supplied curriculum or referenced official guidance; fictional examples are labeled teaching examples.

TEACHABILITY

Pass - facilitator guidance, timing, activities, worksheets, answers, and project criteria are included.

9

Final structured JSON version

Complete machine-readable workshop companion.

The complete structured JSON version of this workshop is delivered as the companion file:

COMPANION FILE

AI_Foundations_for_Founders_Workshop.json

The JSON contains the complete workshop identity, source analysis, materials, learning goals, five modules, ten lessons, final project, visual package, publishing package, QA report, and references.

JSON schema preview

```

{
  "identity": "Workshop metadata and outcomes",
  "source_analysis": "Preserved facts, new curriculum elements, and research basis",
  "materials": "Array of required materials",
  "learning_goals": "Array of six workshop goals",
  "modules": "Array of five module objects",
  "lessons": "Array of ten complete lesson objects",
  "final_project": "Brief, steps, deliverables, checklist, rubric, reflection, challenge",
  "visual_package": "Hero, thumbnail, ten lesson visuals, four diagrams, overview, social",
  "publishing": "SEO, excerpt, headlines, captions, newsletter, video, tags, CTA",
  "qa": "Eight quality checks",
  "references": "Official NIST, FTC, and OpenAI references"
}

```