LUX AUTOMATON

AUTOMATE | INNOVATE | ACCELERATE

AI FOUNDATIONS FOR FOUNDERS

UNDERSTAND AI.
CHOOSE THE RIGHT WORK.
LAUNCH YOUR FIRST PILOT.

A practical, privacy-aware workshop for founders who want useful AI results without losing human control.

LUX AUTOMATON WORKSHOP

Leave with a Founder AI Pilot Blueprint

One real workflow, one tested prompt, clear privacy rules, human approval, and a 30-day plan.

UNDERSTAND

Explain what AI is, where it helps, and where it cannot be trusted without review.

CHOOSE

Rank business opportunities by value, frequency, readiness, risk, and reversibility.

LAUNCH

Scope a controlled first workflow and build a measurable 30-day experiment.

Five modules. Ten lessons. One useful final project.

01 What AI Actually Is	20 min
02 Where Automation Helps First	25 min
03 Private vs. Public AI Tools	25 min
04 Writing Useful Prompts	30 min
05 Choosing Your First Workflow	20 min

Protect the data. Verify the output. Keep the human in charge.

- Use public, synthetic, de-identified, or explicitly approved data during exercises.
- Do not paste passwords, customer records, payment details, or confidential files into public tools.
- Verify AI output against the original source before using it.
- Keep human approval on customer-facing and high-impact decisions.
- Name the person who can pause or stop a pilot.

Explain → Map → Build → Test → Plan

Every activity creates part of the final Founder AI Pilot Blueprint.

1 EXPLAIN
Build the mental model

2 MAP
Find real workflow friction

3 BUILD
Create prompts and controls

4 TEST
Use normal and edge cases

5 PLAN
Launch a 30-day pilot

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 01 • 10 MIN

AI, MACHINE LEARNING, AND GENERATIVE AI

Build a founder-friendly mental model for what AI actually is.

MODULE 1

What AI Actually Is

20 minutes

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

THE FOUNDER'S AI LANDSCAPE

Different capabilities solve different kinds of work.

RULE-BASED AUTOMATION

Follows explicit steps.
Best for stable, repetitive work.

PREDICTIVE AI

Estimates likelihoods or labels.
Best with structured examples.

AI

PATTERN-BASED COMPUTING SYSTEMS

GENERATIVE AI

Creates drafts, summaries, images, code, and variations.

HUMAN JUDGMENT

Sets goals, handles nuance, approves, and remains accountable.

Founder principle **AI is a capability. Workflow design turns it into business value.**

Three capability categories - plus human judgment

RULE-BASED AUTOMATION

Follows explicit steps and conditions. Best for stable, repeatable work.

PREDICTIVE AI

Estimates a label, score, or likelihood from patterns in examples.

GENERATIVE AI

Creates drafts, summaries, images, code, and variations from context.

HUMAN JUDGMENT

Sets goals, handles nuance, approves, and remains accountable.

Input → capability → output → review → action

INPUT
Approved data or trigger

→

CAPABILITY
Rule, prediction, or generation

→

OUTPUT
Draft, score, summary, or action

→

REVIEW
Human check matched to risk

→

ACTION
Send, decide, store, or stop

Every step needs an owner.

Where AI is often useful

- Summarize approved notes
- Transform content for another audience
- Extract fields from consistent documents
- Create first drafts and alternatives
- Classify or route routine inputs
- Surface questions and patterns for review

Where AI can fail

MISSING FACTS

The system may not know your current price, policy, or history.

PLAUSIBLE ERRORS

Fluent language can still contain invented or unsupported details.

INCONSISTENCY

The same prompt may produce different emphasis across runs.

NO ACCOUNTABILITY

AI cannot own the business consequence or final authority.

Confidence is a writing style - not evidence.

Build your Capability Boundary Map

Separate AI assistance from human authority.

- Choose one recurring task.
- Mark what AI may draft, summarize, classify, or suggest.
- Mark what a person must verify, approve, or decide.
- Assign low, moderate, or high consequence.
- Write one explicit no-go boundary.

DELIVERABLE

One-page capability boundary map with task, reviewer, risk level, and no-go rule.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 03 • 12 MIN

FIND REPETITION, FRICTION, AND DELAY

Look for the work that drains time before buying another tool.

MODULE 2

Where Automation Helps First

25 minutes

Look for repetition, friction, and delay

REPEAT

The same data or wording is entered again.

SEARCH

People hunt across folders, inboxes, or tools.

WAIT

Work pauses because ownership or approval is unclear.

REWORK

Output is corrected, reformatted, or rebuilt.

Start with the work - not the tool.

Task, workflow, or decision?

TASK

One unit of work. Example: draft a follow-up.

WORKFLOW

Connected steps, owners, handoffs, and outcomes.

DECISION

A choice that changes what happens next.

AI may assist a task. Automation may move the workflow. A person may own the decision.

Seven questions before automation

1
What triggers the workflow?

2
What inputs are required?

3
What steps happen today?

4
Who owns each step?

5
What output is produced?

6
Where does work wait or repeat?

7
What exceptions break the normal path?

AUTOMATION FIT MATRIX

Match the workflow to the right level of AI involvement.

HIGH

JUDGMENT
AND CONSEQUENCE

HUMAN-LED

High judgment / low repetition
Use AI for research or options only.

AI-ASSISTED

High judgment / high repetition
AI drafts; a person reviews and decides.

MANUAL OR REDESIGN

Low repetition / low judgment
Fix the process before automating.

AUTOMATION CANDIDATE

Low judgment / high repetition
Automate after tests and exception rules.

LOW
LOW

HIGH

REPETITION AND STANDARDIZATION →

The Lux Opportunity Scorecard

VALUE
Business benefit or risk reduction

FREQUENCY
How often the workflow occurs

STANDARDIZATION
How repeatable the work is

DATA READINESS
Approved inputs are available

RISK
Cost or harm of a wrong output

REVERSIBILITY
How easily the workflow can be corrected or stopped

A strong first pilot is useful, measurable, reversible, and close to human review.

Choose the right first win - not the biggest promise

WEEKLY CLIENT UPDATE

Frequent, source notes available, easy to review, easy to correct.

FINAL CONTRACT REVIEW

High consequence, complex exceptions, requires qualified authority.

RESEARCH SUMMARY

Useful and testable, but sources and currency must be checked.

Pilot choice: weekly client update.

Rank three AI opportunities

Use evidence. Challenge one assumption.

- Score value, frequency, standardization, data readiness, risk, and reversibility.
- Mark whether a knowledgeable reviewer is close to the output.
- Mark whether shadow testing is possible.
- Rank the three opportunities.
- Write why the top choice is first - and why the others wait.

DELIVERABLE

A scored shortlist and one recommended first candidate.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 05 • 12 MIN

PUBLIC, PRIVATE, AND ENTERPRISE AI

Understand how tool type changes data handling, control, and accountability.

MODULE 3

Private vs. Public AI Tools

25 minutes

PUBLIC VS. PRIVATE AI DECISION TREE

Classify the information before choosing the tool.

WHAT INFORMATION WILL ENTER THE TOOL?

Does it include customer data, private business data, regulated information, credentials, or unreleased strategy?

NO

YES

NO

Use an approved public tool if company policy allows it.
Still remove unnecessary identifiers and verify the output.

YES / UNSURE

Use an approved private or enterprise environment, or do not use AI until the data owner approves.

Every path ends with: human review, permissions, and a record of the decision.

Three tool environments - different controls

PUBLIC CONSUMER

Standard terms. Limited organization administration. Best reserved for public, generic, synthetic, or carefully de-identified information when policy permits.

BUSINESS / ENTERPRISE

May add contractual data commitments, admin controls, identity, retention options, and audit features.

PRIVATE / DEDICATED

Dedicated service, controlled cloud, or self-managed environment. Controls must still be verified.

Do not rely on labels. Verify the actual terms and configuration.

Seven questions before sensitive use

1. Is content used to train shared models?
2. How long are prompts, files, and outputs retained?
3. Who can access the data?
4. Can admins manage accounts and permissions?
5. Is data encrypted in transit and at rest?
6. Are audit logs and deletion controls available?
7. What contract and security commitments apply?

Classify data before you prompt

PUBLIC

Approved for anyone to see

INTERNAL

Operational but not public

CONFIDENTIAL

Could harm a person, customer, or business if disclosed

RESTRICTED

Credentials, regulated records, highly sensitive personal or proprietary data

Four controls after classification

MINIMIZE

Use only what the task needs.

DE-IDENTIFY

Remove names and unique details when possible.

AUTHORIZE

Confirm the user and tool are approved.

RETAIN

Know where prompts, files, outputs, and logs live - and when they are deleted.

A summary of confidential information is still confidential.

Build your AI Data Rule Card

GREEN - APPROVED

Public or synthetic information in an approved tool.

YELLOW

Controlled: internal or confidential use only in an approved environment with review.

RED

Do not enter: credentials, secrets, restricted records, or unapproved sensitive data.

Add an escalation contact for uncertain cases.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 07 • 15 MIN

WRITE USEFUL PROMPTS WITH CRAFT

Give AI the context, assignment, format, and quality test it needs.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

THE LUX CRAFT PROMPT FRAMEWORK

Five elements that turn a vague request into a usable business instruction.

C
CONTEXT
What is happening? Who is the audience? What source is trusted?

R
ROLE
What expertise or perspective should the AI apply?

A
ASSIGNMENT
What exact task and business deliverable must be produced?

F
FORMAT + BOUNDARIES
What structure, tone, length, and exclusions are required?

T
TEST
How will quality be checked? What should AI flag or ask?

Prompting principle **Be clear enough to guide the work, then verify against real standards.**

CRAFT is a compact operating instruction

C CONTEXT
Situation, audience, source, goal

R ROLE
Useful perspective - not authority

A ASSIGNMENT
Exact action and deliverable

F FORMAT + BOUNDARIES
Structure, tone, length, exclusions

T TEST
Quality criteria and missing-info rule

Vague request vs. business instruction

WEAK

"Write a follow-up email."

CRAFT

Using the approved call notes below, draft a 150-word follow-up for a prospective bookkeeping client. Confirm the three stated needs, propose a discovery call, and use a warm professional tone. Do not invent prices, dates, guarantees, or commitments. List missing information before the draft. Check every factual statement against the notes.

Draft → Inspect → Refine → Verify → Save

DRAFT
Run the first CRAFT prompt

→

INSPECT
Accuracy, completeness, usefulness, tone, risk

→

REFINE
Change one important variable

→

VERIFY
Compare with source and checklist

→

SAVE
Version, owner, data class, limits

Three cases every reusable prompt should face

NORMAL CASE

Complete approved information. Does the prompt produce the required output?

MISSING INFORMATION

Does the prompt ask, flag, or insert a visible placeholder instead of guessing?

EXCEPTION CASE

Does the workflow route unusual or conflicting cases to a person?

Build and test one reusable CRAFT prompt

- Write the five CRAFT elements.
- Run or simulate one normal case.
- Test one missing-information case.
- Test one exception case.
- Revise one instruction at a time.
- Save version, owner, approved data class, reviewer, and limitations.

DELIVERABLE

Version-two prompt, three-case test log, verifier checklist, and approval record.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 09 • 10 MIN

CHOOSE A LOW-RISK FIRST WORKFLOW

Start narrow, reversible, measurable, and close to human review.

MODULE 5

Choosing Your First Workflow

20 minutes

First Workflow Canvas

TRIGGER
What starts it?

APPROVED INPUT
What may enter?

AI ACTION
What bounded task?

OUTPUT
What one deliverable?

HUMAN REVIEW
Who checks?

DESTINATION
Where after approval?

EXCEPTIONS
What routes to a person?

OUT OF SCOPE
What will not happen?

Narrow, reversible, measurable, and close to review

- One trigger
- One approved input set
- One primary output
- One named reviewer
- Explicit in-scope and out-of-scope behavior
- Exception routing
- Clear stop authority

PILOT HYPOTHESIS

If we use an approved AI prompt to draft follow-ups from approved call notes, then account leads can prepare accurate responses faster without increasing corrections or complaints.

LUX AUTOMATON
AUTOMATE | INNOVATE | ACCELERATE

LESSON 10 • 10 MIN

BUILD YOUR 30-DAY AI PILOT ROADMAP

Set owners, checkpoints, metrics, guardrails, and a clear go-or-stop decision.

LUX AUTOMATON WORKSHOP

Five evidence-building phases

DAYS 1-5
Baseline + controls

DAYS 6-10
Build + internal test

DAYS 11-15
Shadow comparison

DAYS 16-25
Controlled use

DAYS 26-30
Evidence review

End decision: GO • REVISE • STOP

Speed alone is not success

BALANCED MEASURES

Cycle time
Reviewer effort
Acceptance without major edits
Rework and exception rate
Staff or customer feedback
Incidents

STOP CONDITIONS

Unauthorized data use
Unsupported claims
Missed mandatory steps
Material customer harm
No audit trail
Failure above tolerance

Founder AI Pilot Blueprint

- Current-state workflow map
- Scored opportunity shortlist
- Tool class and data-use decision
- AI data rule card
- Version-two CRAFT prompt
- Three-case test log and verifier checklist
- First Workflow Canvas
- 30-day roadmap and go/revise/stop criteria

COMPLETION STANDARD

20-point rubric. Pilot-ready = 16+ with no criterion below 3.

Five criteria • four points each

1. Workflow fit and scope — 4 points
2. Privacy, tool choice, and governance — 4 points
3. Prompt and test quality — 4 points
4. Human oversight and risk controls — 4 points
5. 30-day implementation and measurement — 4 points

LUX AUTOMATON

AUTOMATE | INNOVATE | ACCELERATE

AI FOUNDATIONS FOR FOUNDERS

UNDERSTAND AI.
CHOOSE THE RIGHT WORK.
LAUNCH YOUR FIRST PILOT.

A practical, privacy-aware workshop for founders who want useful AI results without losing human control.

BRING ONE RECURRING BUSINESS FRUSTRATION.

LEAVE WITH ONE RESPONSIBLE AI WORKFLOW READY TO TEST.